

Storytelling with Data

Module 5: Analyze before you speak—audience analysis

Scott Spencer
Faculty and Lecturer
Columbia University

Agenda

Upcoming deliverable

Today's objectives

Audience: questions, biases

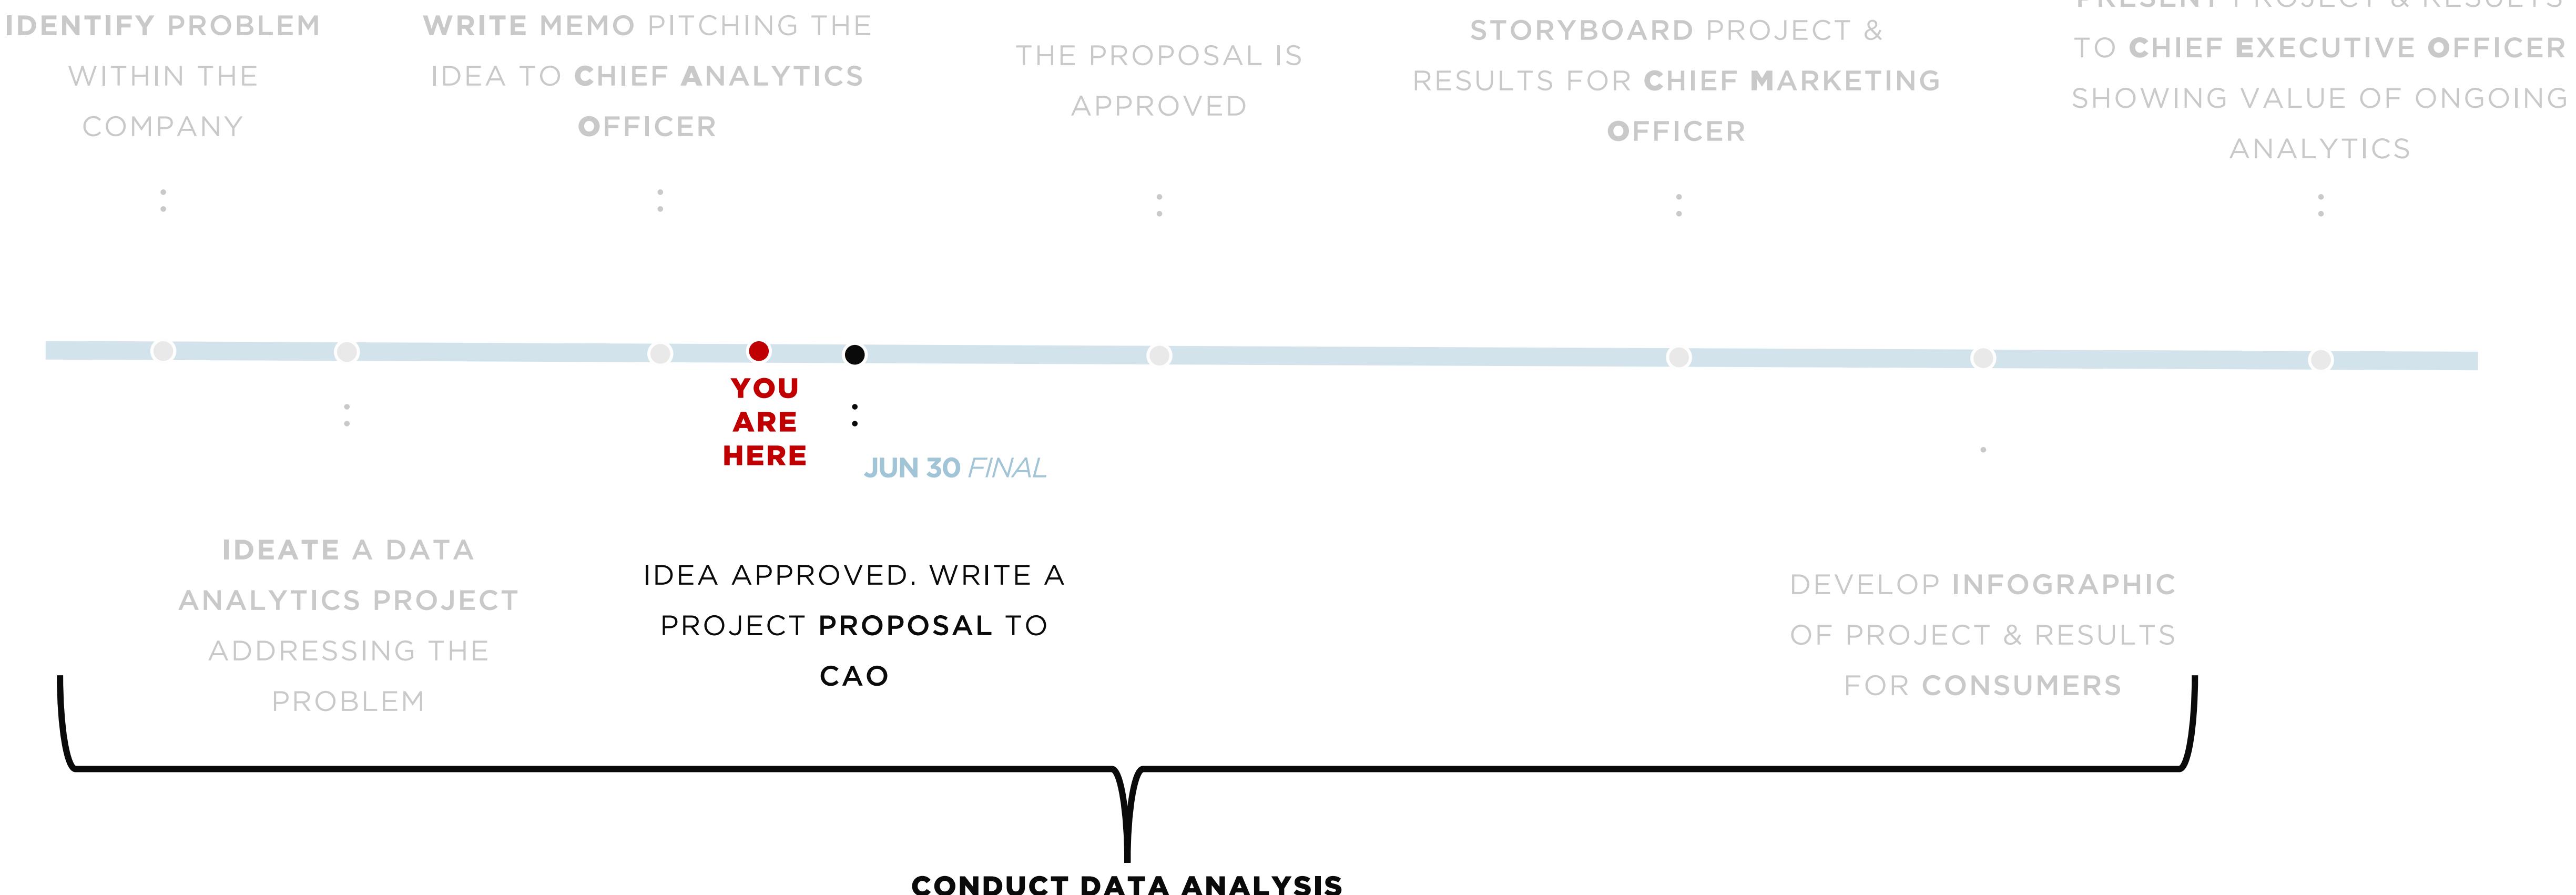
Audience: cultural cues, emotional display

Fair, influential communication

Upcoming deliverable

Upcoming deliverable

750-word brief proposal – Write a brief proposal to **CAO** detailing your proposed analytics project. Consider background context, problem, data, solution, and impact. At this point you should have data to start an analysis.



Turn generic informational sections,

1

Introduction

Describe context relevant to the proposal.

2

Challenge and Opportunity

Offer necessary details of the problem and identify your solution.

3

Rationale

Explain your solution, how the ideas can be applied to the particular need, and why this solution will work.

4

Cost Analysis

Consider cost, resources required, and expected return on investment, using quantitative or qualitative language.

5

Assessment

Describe how the project will be evaluated both immediately and over time.

6

Conclusion

Include a clear and complete call to action.

into specific messages of your project.

Example brief proposal for a data analytics project. Constraint—750 words or less in main body.

Proposal for exploring game decisions informed by expectations of joint probability distributions

To: Scott Powers, Senior Baseball Analyst, Los Angeles Dodgers
From: Scott Spencer, Faculty and Lecturer, Columbia University

14 February 2019

Our game decisions based on current modeling do not maximize spend per win. We witnessed the mid-market Astros use analytics to overtake us in the 2017 World Series (Luhnow 2018a,b). Our efforts also do not maximize expected wins. But we can. To do so, we need to jointly model probabilities of all game events and base decisions on expectations of those distributions. With adequate computing emerging, we can be first using the probabilistic programming language Stan and parallel processing. To demonstrate the concept, consider a probability model for decisions to steal second base, below, which suggests teams are too conservative, leaving wins unclaimed. This model allows us to ask, for example—should Sanchez steal against Sabathia? Or against Pineda?

1 Our current analyses do not optimize expected wins

Seven terabytes of uncompressed data generated per game overshadow the lack of situational data needed for decision-making that maximizes expected utility. Consider that pitchers, on average, only face 10 percent of major league batters regardless of game state; the reverse is true, too. Or when deciding whether a base runner should attempt to steal against a specific pitcher and catcher in a state of play, say, we are lucky to have any data. Common analyses and heuristics for these situations are inadequate: they not only overfit the data (if any exist), but also offer no manner of estimating changes in probabilities for maximizing expected utility (winning the game).

Accurately quantifying probabilities, and changes thereof, in a given context enable us to answer counterfactuals, from which we can build strategies that maximize our objectives (Parmigiani 2002). This approach is possible at scale using Stan (Carpenter et al. 2017). It's time to jointly model probabilities of all events.

2 Modeling probabilities for steal success illustrates a broader benefit

To see the potential of implementing probability models, let's consider, again, the decision to steal bases, given a specific counterfactual:

PROPOSAL FOR EXPLORING GAME DECISIONS INFORMED BY EXPECTATIONS OF JOINT PROBABILITY DISTRIBUTIONS 2

In a game against New York Yankees, should Milwaukee Brewer's Lorenzo Cain attempt to steal second base with no one else on base and two outs before the seventh inning, against Gary Sanchez as catcher and Michael Pineda as pitcher? What if against Sanchez and CC Sabathia as pitcher?

More specifically, how can we know the expectation that Cain's attempt in each situation increases the probability of expected runs that inning and by how much? Using Stan, I've coded a generative model that along with play outcomes considers various information (runner foot-speed, catcher pop-time) and player characteristics, like pitcher handedness. With the model, we have an answer that also shows the uncertainty. Given 2017 data, this model suggests Cain should steal against Pineda, not Sabathia:

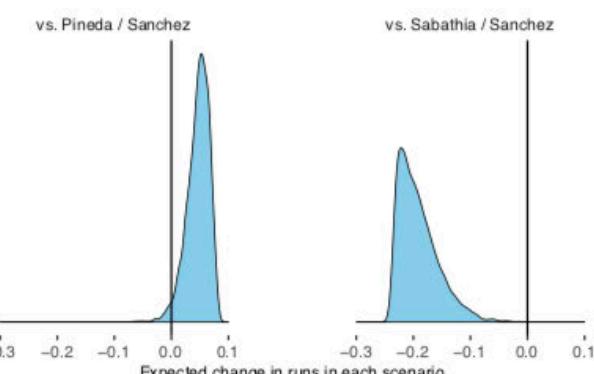


Figure 1. Of the two scenarios, Cain should only attempt to steal against the Sanchez-Pineda duo.

Notably, we get these expectations without multiple trials of either scenario. More generally, this model suggests that on average team managers are too conservative, leaving runs unrealized:

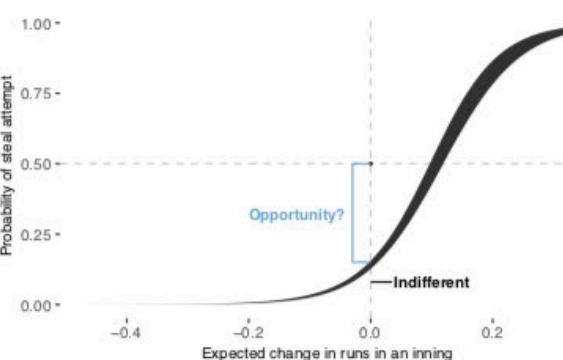


Figure 2. When the change in expected runs is zero, managers should be indifferent to attempted steals, saying go half the time.

The black band represents the range of variation across managers' decisions. At the intersection of indifference, managers tend to say steal only 10 percent of the time, leaving opportunity.

The above is but one example of a more general approach that weighs probabilities of all possible outcomes to maximize expected utility. With broad implementation—jointly modeling the conditional probabilities of all relevant events—we can optimize decisions.

PROPOSAL FOR EXPLORING GAME DECISIONS INFORMED BY EXPECTATIONS OF JOINT PROBABILITY DISTRIBUTIONS 3

3 For value, compare an investment to free-agent costs

A fully-realized model will require significant effort from a team with deep experience in baseball, generative modeling, and Stan. To get the talent, we should compare cost to acquiring expected wins from free-agents. Each win above a replacement-level player costs about 10 million per year (Swartz 2017). As with free-agent value over replacement player, game-time decisions informed from more accurate probabilities should add wins over a season. The scope of what we can answer, moreover, goes beyond in-game strategy (player acquisitions, salary arbitration). More immediately, however, we can begin to implement this approach for specific events, with a scope closer to the example above, being mindful that information learnt are conditional upon unmodeled context.

4 For accuracy, compare model results to betting market odds

Measuring performance of a fully-realized model may seem tricky: we only see the outcome of our decisions. But we can, say, compare the accuracy of our estimates against the betting market where interested investors are trying to forecast game outcomes.

5 Conclusion

The mid-market Astros show teams can do more with information. Millions in additional revenue—and more wins—await discovery through a joint, probability model of all events from which we can maximize conditional expectations. Let's discuss how to draw the talent for a title worth our spend.

6 References

- Carpenter, Bob, et al. 2017. "Stan: A Probabilistic Programming Language." *Journal of Statistical Software* 76 (1): 1–32.
Luhnow, Jeff. 2018a. "How the Houston Astros are winning through advanced analytics." *McKinsey Quarterly* 13 June 2018: 1–9.
———. 2018b. "A view from the front lines of baseball's data-analytics revolution." *McKinsey Quarterly* 5 July 2018: 1–8.
Parmigiani, G. 2002. "Decision Theory: Bayesian." In *International Encyclopedia of the Social Behavioral Sciences*, 3327–34.
Swartz, Matt. 2017. "The Recent History of Free-Agent Pricing." <https://www.fangraphs.com/blogs/the-recent-history-of-free-agent-pricing/>.

Readability Statistics	
Counts	
Words	720
Characters	3,997
Paragraphs	16
Sentences	35
Averages	
Sentences per Paragraph	4.3
Words per Sentence	18.1
Characters per Word	5.3
Readability	
Flesch Reading Ease	33.2
Flesch-Kincaid Grade Level	13
Passive Sentences	0%

Today's Objectives

Objectives

1

Articulate the need for **audience analysis** and sensitivity in the applied analytics setting.

2

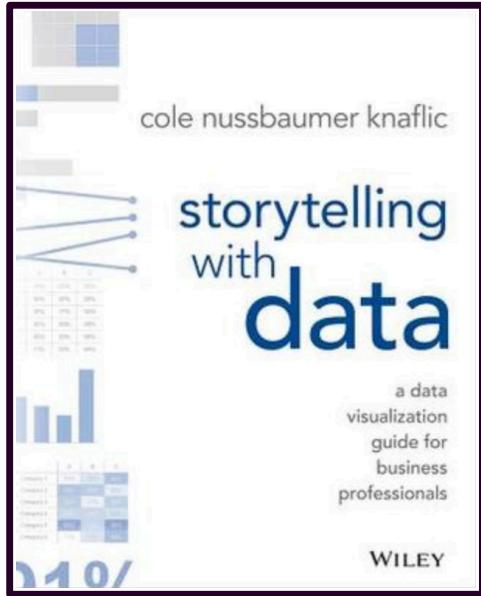
Consider **professional, demographic, cultural, and personal issues** when building, contributing to, or managing an analytics team or project.

Audience analysis and biases

Storytelling with data

Knafllic

The author is a consultant focused on visual displays. Her experience arose from human resources in Google where she applied theory learned as a student of Yale's Edward Tufte.



Audience: questions to specifically answer

Who is the audience or decision maker?

What is your relationship with them?

What do you need them to know or do?

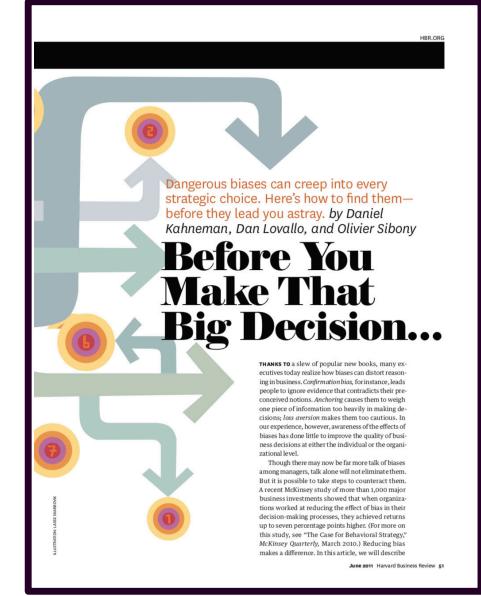
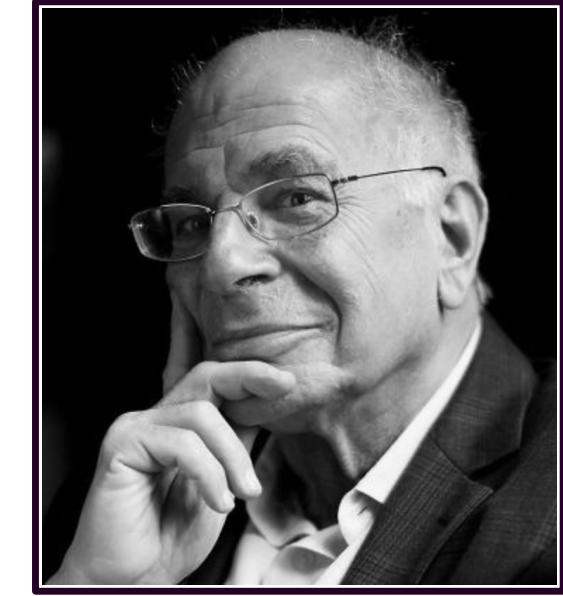
How will you communicate with them?

What tone do you want your communication to set?

What data are available to help make your point?

Is your audience familiar with these data?

What are audience biases as related to your messages?



Before you make that big decision...

Kahneman, co-authors

Awarded the Nobel-Prize in economics and senior scholar at Princeton, Kahneman introduced the idea of cognitive biases, and their impact on decision making.

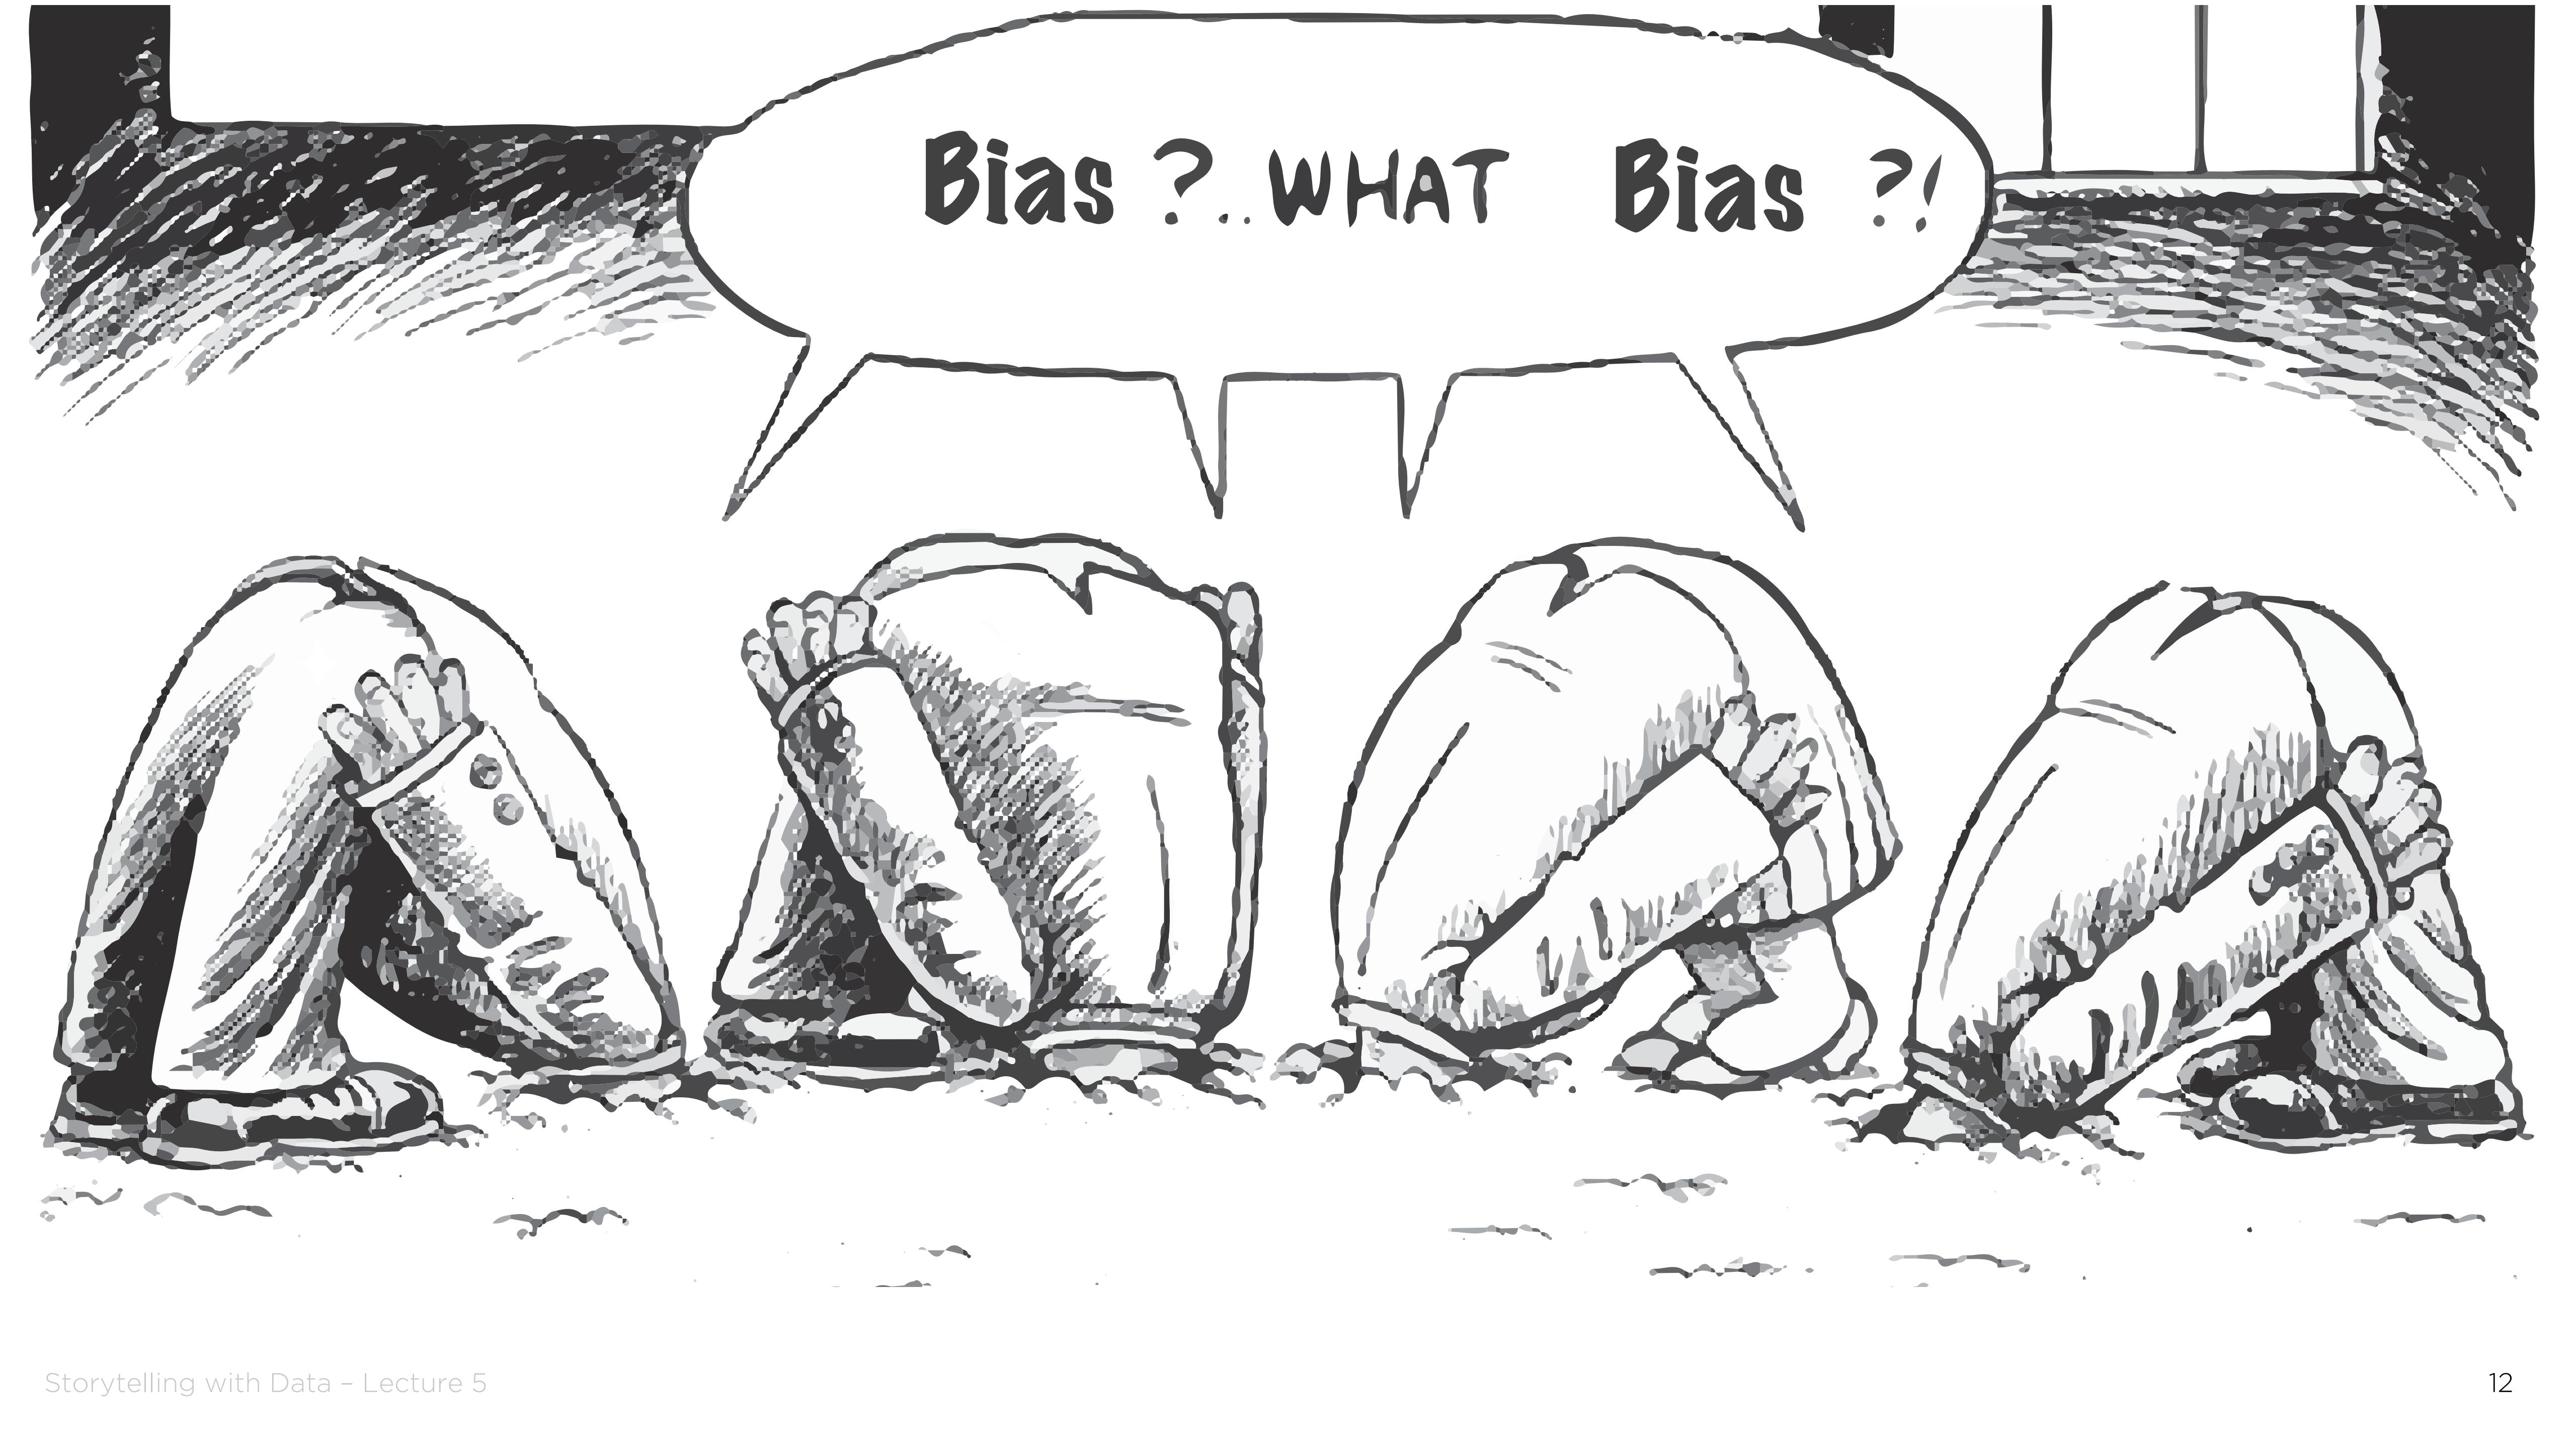
Two modes of thinking

Intuitive (system one) thinking, impressions, associations, feelings, intentions, and preparations for actions flow effortlessly. This system mostly determines our thoughts. System one uses **heuristics**, has **biases**.

Reflective (system two) thinking is slow, effortful, and deliberate.

Both are continuous, but system two typically monitors things, and only steps in when stakes are high, we detect an obvious error, or rule-based reasoning is required.

It's very hard to remain aware of our own biases, so we need to **develop processes** that identify them and, most importantly, get **feedback from others** to help protect against them.



Bias ?..WHAT Bias ?!

1

Check for self-interested biases

Is there any reason to suspect the team making the recommendation of errors motivated by self-interest?

Review the proposal with extra care, especially for over optimism.

2

Check for the affect heuristic

Has the team fallen in love with its proposal?

Rigorously apply all the quality controls on the checklist.

3

Check for groupthink

Were there dissenting opinions within the team? Were they explored adequately?

Solicit dissenting views, discreetly if necessary.

4

Check for saliency bias

Could the diagnosis be overly influenced by an analogy to a memorable success?

Ask for more analogies, and rigorously analyze their similarity to the current situation.



Check for confirmation bias

Are credible alternatives included along with the recommendation?

Request additional options.



Check for availability bias

If you had to make this decision in a year's time, what information would you want, and can you get more of it now?

Use checklists of the data needed for each kind of decision.



Check for anchoring bias

Where are the numbers from? Can there be ... unsubstantiated numbers? ... extrapolation from history? ... a motivation to use a certain anchor?

Re-anchor with data generated by other models or benchmarks, and request a new analysis.



Check for halo effect

Is the team assuming that a person, organization, or approach that is successful in one area will be just as successful in another?

Eliminate false inferences, and ask the team to seek additional comparable examples.

9

Check for sunk-cost fallacy, endowment effect

Are the recommenders overly attached to past decisions?

Consider the issue as if you are a new executive.

10

Check for overconfidence, optimistic biases, competitor neglect

Is the base case overly optimistic?

Have a team build a case taking an outside view: use war games.

11

Check for disaster neglect

Is the worst case bad enough?

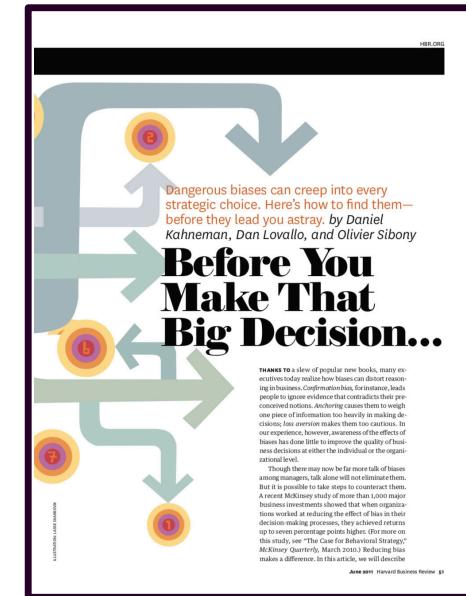
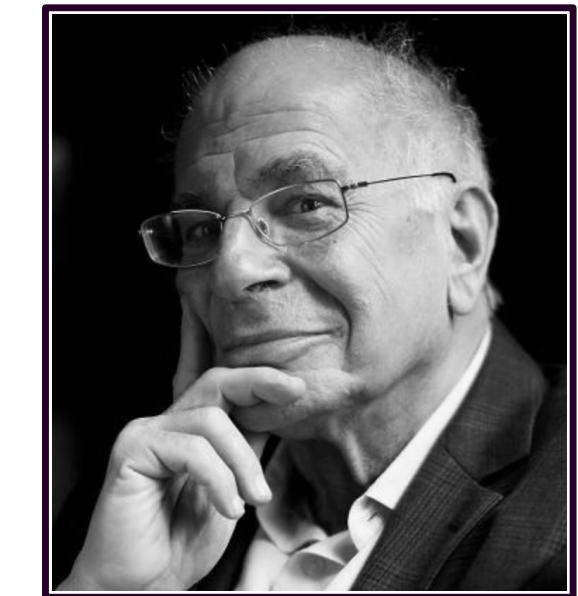
Have the team conduct a premortem: imaging that the worst has happened, and develop a story about the causes.

12

Check for loss aversion

Is the recommending team overly cautious?

Align incentives to share responsibility for the risk or to remove risk.



Before you make that big decision...

Kahneman, co-authors

Awarded the Nobel-Prize in economics and senior scholar at Princeton, Kahneman introduced the idea of cognitive biases, and their impact on decision making.

Keeping out the appearance of bias

Present ideas from a **neutral perspective**. Becoming too emotional suggests bias.

Make **analogies and examples comparable** to the proposal.

Genuinely **admit uncertainty** in the proposal, and **recognize multiple options**.

Identify **additional data** that may provide new insight.

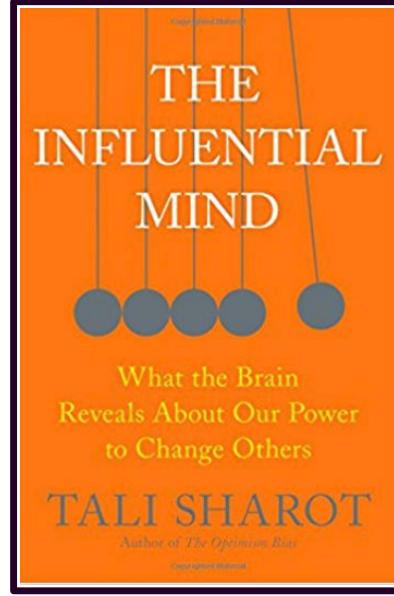
Consider **multiple anchors** in the proposal.

Fair, influential communication

The influential mind

Sharot

A London neuroscientist, her research focuses on decision-making, emotion, and influence.



Learn what your audience is thinking

If we want to affect the behaviors and beliefs of the person in front of us, we need to first understand what goes on inside their head.

Formula for changing beliefs

Four factors come into play when we form a new belief: our **old belief** (this is technically known as the “prior”), our confidence in that old belief, the new evidence, and our confidence in that evidence.

Find common ground with audience's beliefs

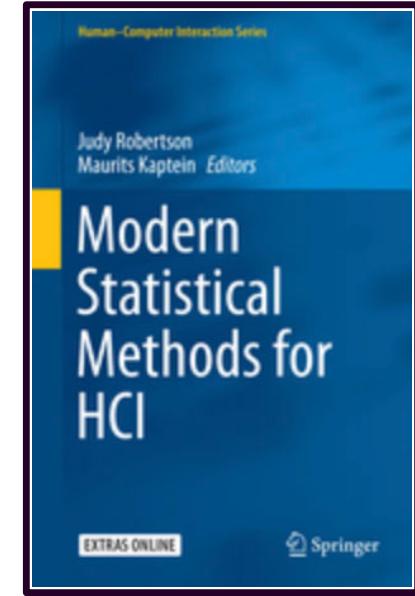
When you provide someone with new data, they quickly accept evidence that confirms their preconceived notions (what are known as prior beliefs) and assess counterevidence with a critical eye.

Focusing on what you and your audience have in common, rather than what you disagree about, enables change.

Ch. 13, Fair statistical communication in HCI

Dragicevic

He is a researcher, focusing on psychology of data visualization for judgment and decision making, and on transparent statistical communication.



Report effect sizes, interval estimates, interpretations

Whether exploring or confirming analyses, show results using an *estimation* approach—use graphs to show **effect sizes** and **interval estimates**, and offer **nuanced interpretations** of results. *Avoid the pitfalls* of dichotomous tests and p-values.

Tip 19: be creative, use visuals

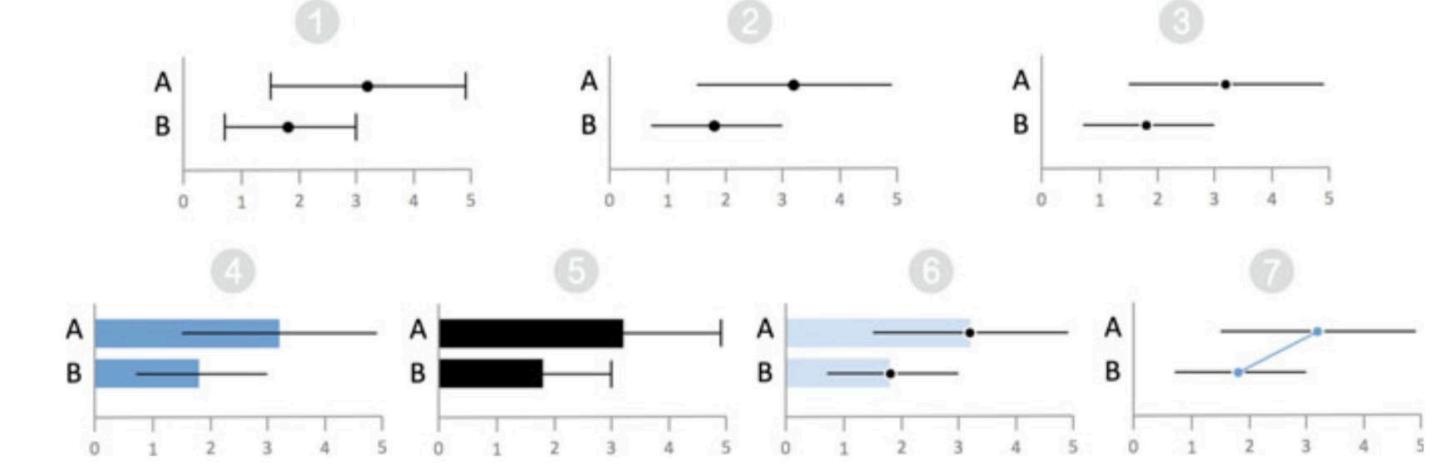
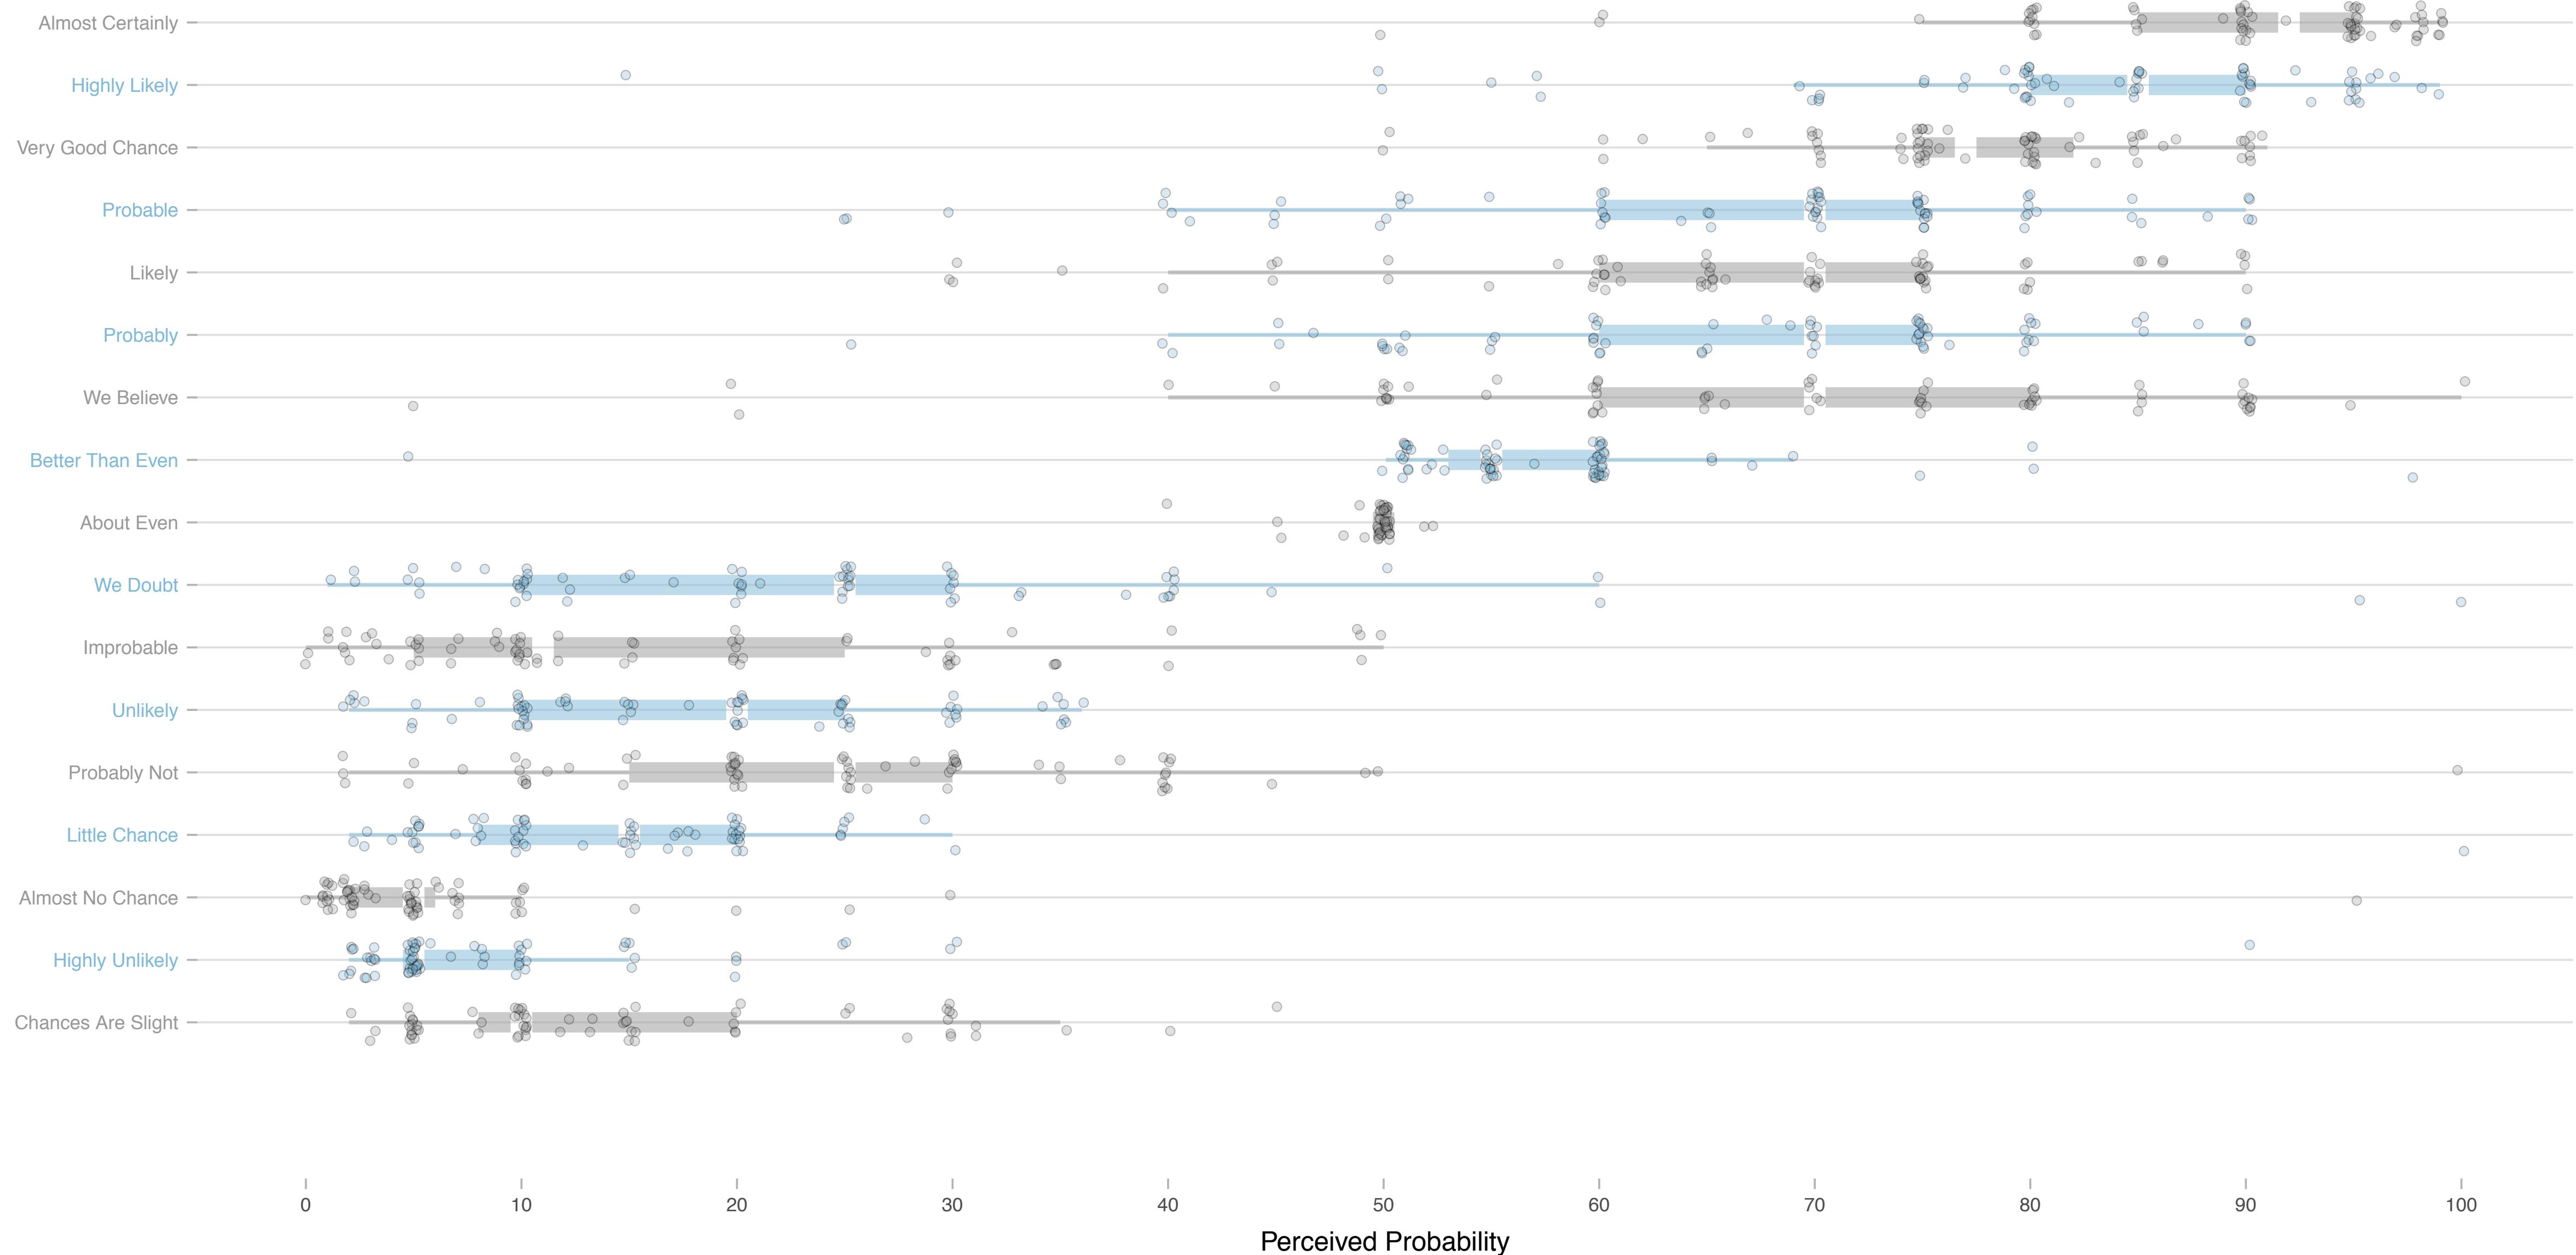


Fig. 13.9 Seven ways of plotting effect sizes with confidence intervals

Tip 26: honestly convey uncertainty

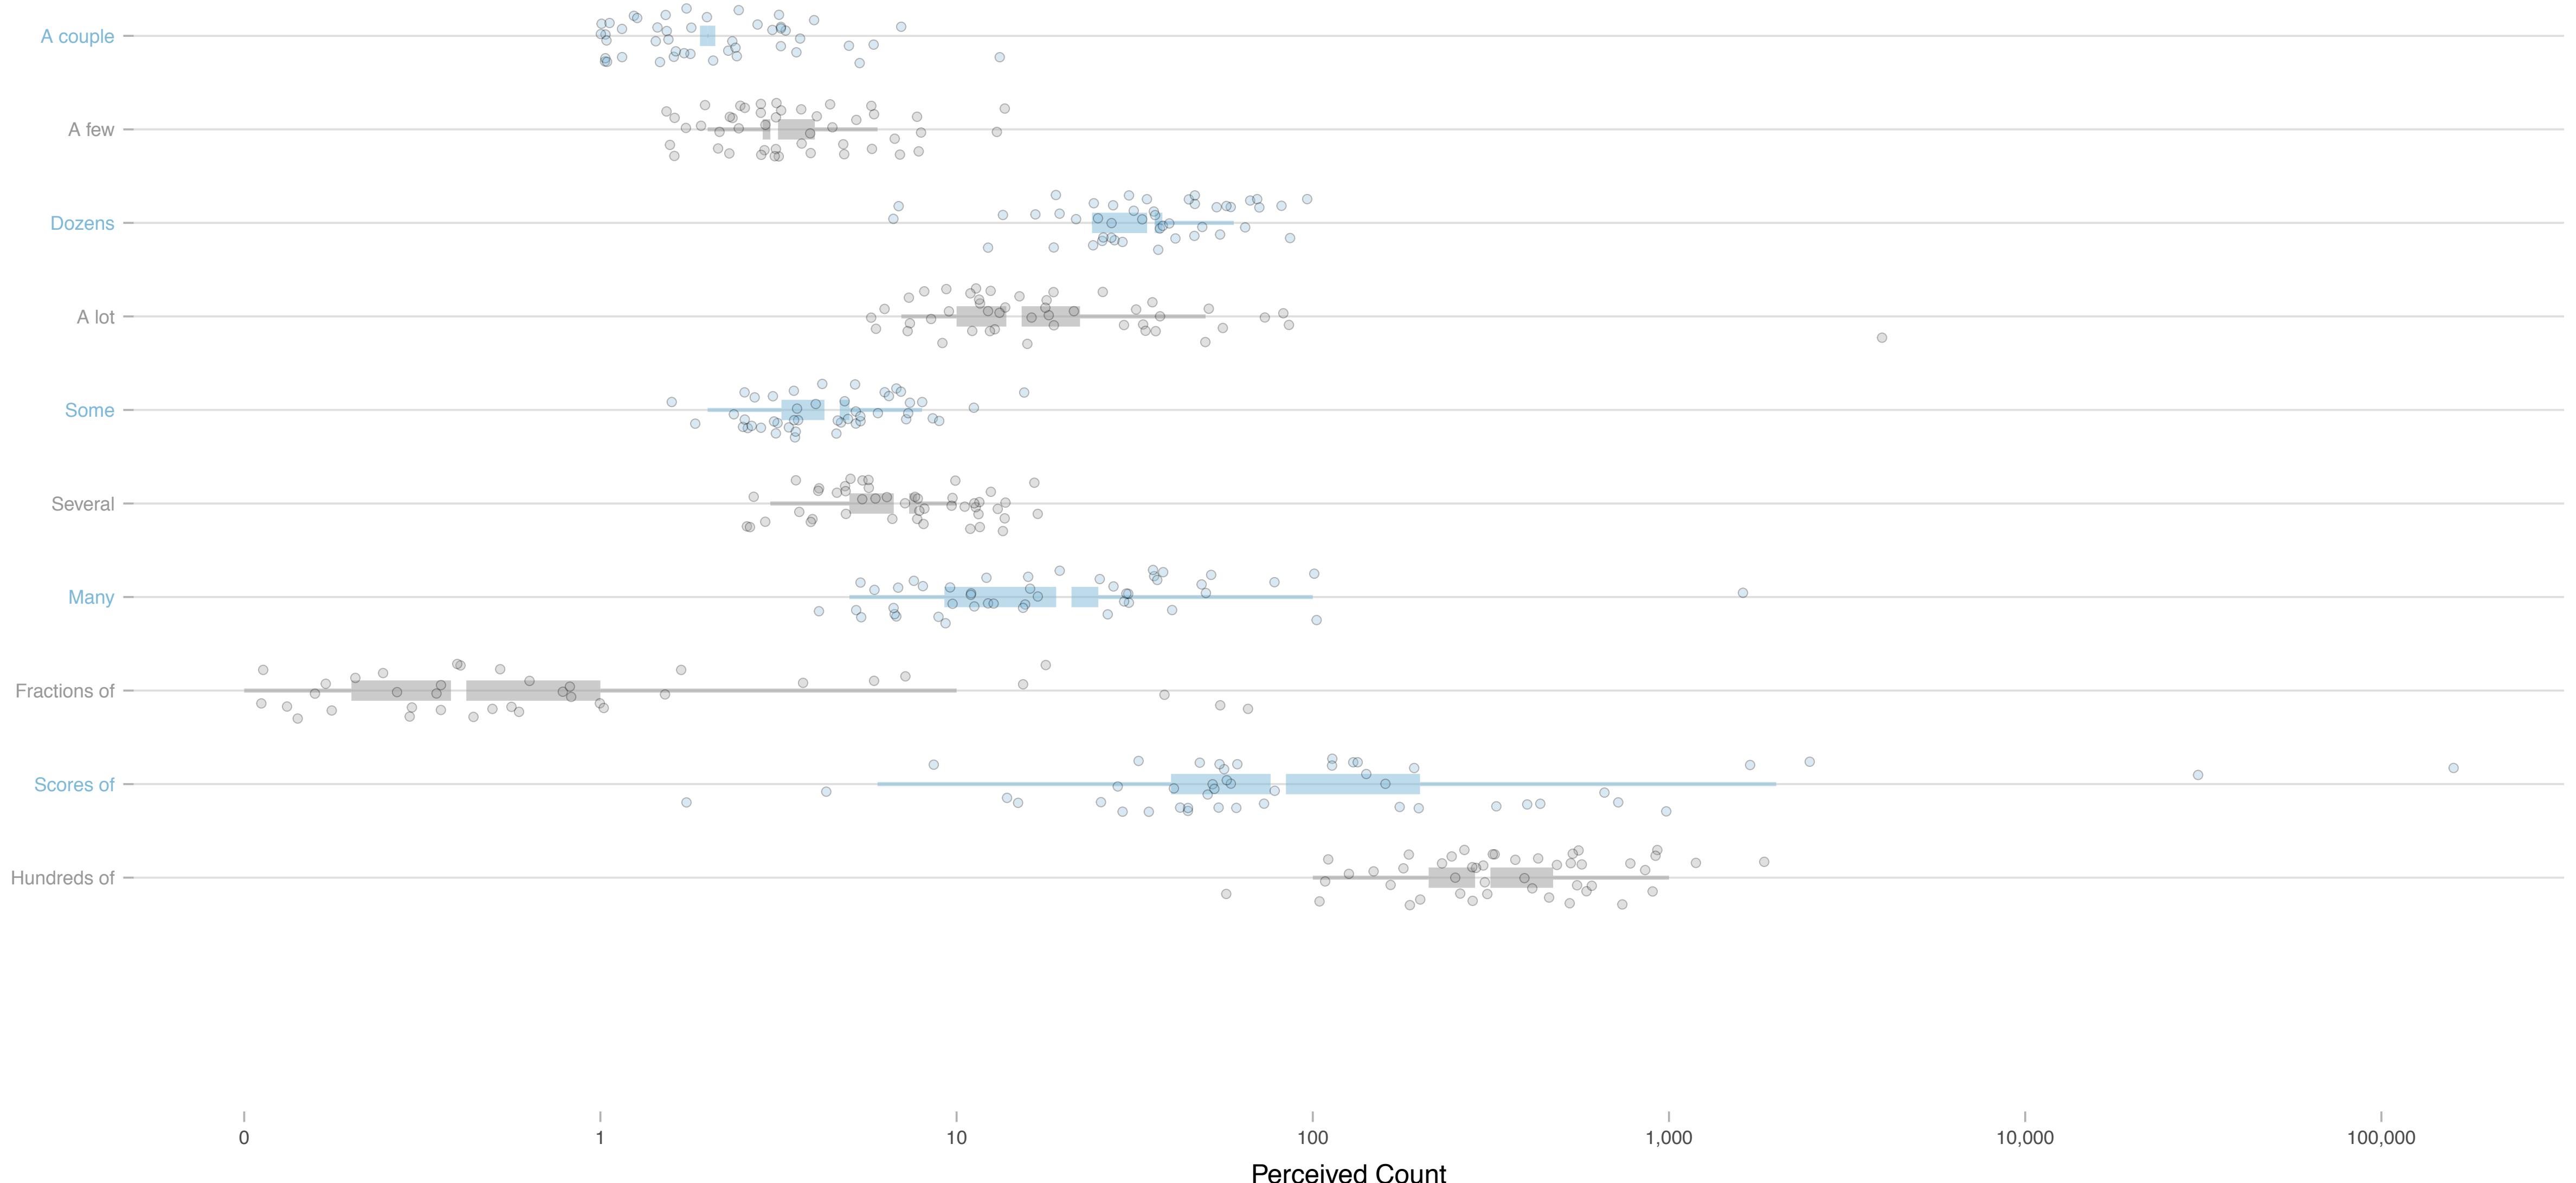
The use of vague language is necessary for **acknowledging and honestly conveying the uncertainty present in effect size estimates**. Vague language — which is not the same as ambiguous language — plays a key role in reasoning.

Language describing probability are imprecise, depend upon audience and context



Spencer, Scott. *Quantitative persuasion amid uncertainty*. Forthcoming. Print.
Barclay, Scott et al. *Handbook for Decision Analysis*. Decisions and Designs, Inc., 1977. Print.
zonization. "Perceptions of Probability and Numbers." github.com/githhub.com, 18 Aug. 2015. Web. 26 Dec. 2018.

Language describing quantities are imprecise, depend upon audience and context



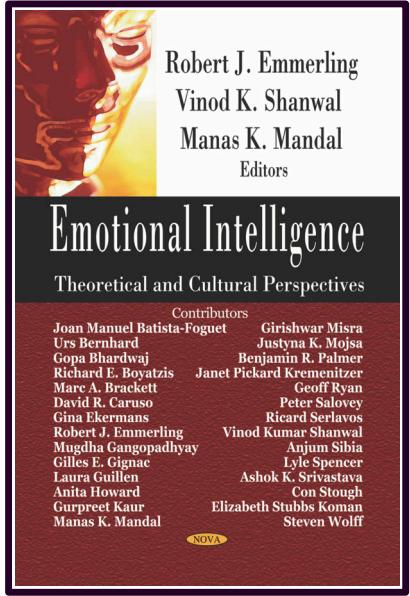
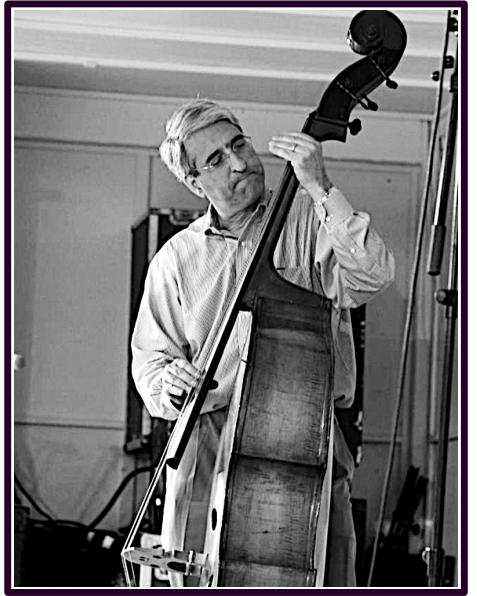
Spencer, Scott. *Quantitative persuasion amid uncertainty*. Forthcoming. Print.
Barclay, Scott et al. *Handbook for Decision Analysis*. Decisions and Designs, Inc., 1977. Print.
zonination. “Perceptions of Probability and Numbers.” github.com/18 Aug. 2015. Web. 26 Dec. 2018.

Cultural cues, emotional display

Emotional intelligence theoretical and cultural perspectives

Salovey, co-authors

Yale president (and double bassist), Salovey introduced the idea of emotional intelligence. This book collects numerous authors' perspectives, examining the **cross-cultural similarity of the concept of emotional intelligence**.



**Emotions are universal,
but expression
depends on culture**

**Emotional Intelligence
guides communication**

**Higher and lower
context cultures**

Basic emotions are **perceived similarly** across the world, but display rules for emotional **expression vary** from culture to culture. That is, the norms pertaining to how—and to what intensity—certain emotions should be expressed within social contexts.

“Emotional intelligence involves the ability to monitor one’s own and **others’ feelings and emotions**, to discriminate among them, and to use this information to **guide one’s thinking and actions**” including “our **ability to communicate effectively**.”

Caveat: much of the research in this field has originated from Western cultures. This reference includes other perspectives.

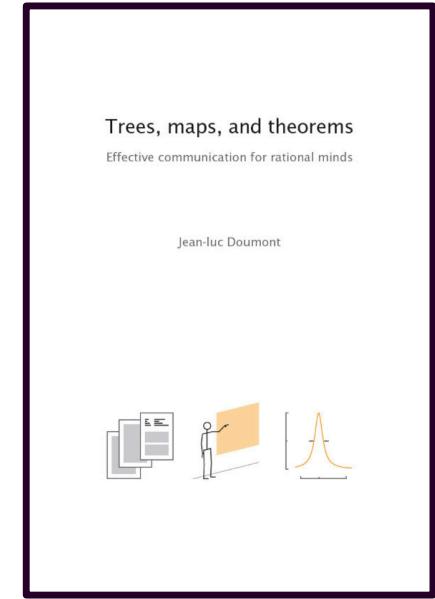
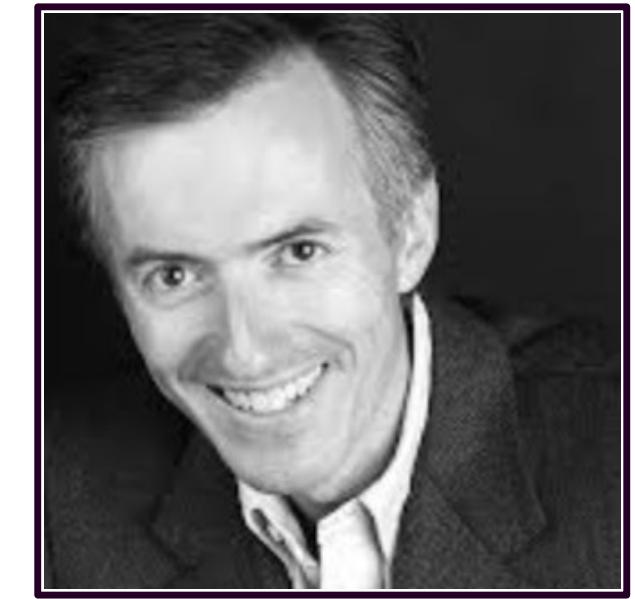
In **high context cultures**, unlike the US, the internal meaning of a message is usually embedded deeper in the information; **not everything is explicitly written**.

Questions for discussion

How have your experiences compared with or differed from what we have reviewed so far?

If you were the decision maker for approval of a data analytics project, what type of information or what form of communication would you want to have?

Group help



Trees, maps, and theorems

Doumont

An engineer from the Louvain School of Engineering and PhD in applied physics from Stanford University, Jean-luc Doumont wrote this book to help engineers, scientists, and managers with business communication.

Determine purpose(s) of the review

When reviewing someone else's document, center yourself on the **purpose that was agreed upon**, such as clarity, accuracy, or correctness.

Should this purpose be multiple, **review one aspect at a time, focusing on content first**.

Reasoning before typos

Help, don't judge

Structure the review

Typos are usually more conspicuous than reasoning flaws, but also less important.

In your comments to the authors, **strive to help, not to judge**.

First, provide a **global assessment**, to place further comments in proper perspective. As a rule, point out the **weaknesses**, to prompt improvements, but also the **strengths**, to increase the authors' willingness to revise the document and to learn.

Groupwork, feedback from peers

Pair up, share your proposal summary, and practice giving feedback on your drafts based on the requirements just discussed.

Going forward

For Next Week, Module 6:

Agenda next week

The minimum

Next deliverable, final 750-word (or less) proposal
The storytelling process
As you review the material, think about how you might storyboard *The Next Rembrandt*.

Riche, Nathalie Henry et al. *Data-Driven Storytelling*. CRC Press, 2018. Print. Selected readings.

Read for more perspective on the process of data-driven, visual stories.

Lee, Bongshin et al. “More Than Telling a Story: Transforming Data Into Visually Shared Stories.” IEEE Computer Graphics and Applications 35.5 (2015): 84–90. Web.

Read for more perspective on the process of data-driven, visual stories.

McCloud, Scott. *Understanding Comics: the Invisible Art*. Kitchen Sink Press, 1993. Print. Selected readings.

Read for ideas on framing, and think about how it relates to storyboarding.

NAPA Cards. *Narrative Patterns for Data-Driven Storytelling*. Web. <http://napa-cards.net>

Compare differences in narrative patterns.

Q's in threes,

**Of metaphor
and story**

**Too much,
too little?**

**Help on
next steps?**

Thinking back to last week's material of metaphor, what metaphors may help you to describe or explain your project and results? Would the metaphor be equally effective with CAO as with CMO?

In the past, I've received feedback that the course is quite challenging. As a group, your draft briefs were well done! Tell me where you feel overwhelmed on specific topics, and I'll try to provide more guidance.

How do you perceive your progress so far, given our timeline? As you get stuck, reach out to both your team (this class) and I for guidance.

**See you
next week!**

